

ALEX MORGAN

San Francisco, CA | alex.morgan@email.com | github.com/alexmorgan | linkedin.com/in/alexmorgan

PROFESSIONAL SUMMARY

Full Stack AI Engineer with 2+ years of experience building high-performance distributed systems, autonomous LLM agents, and scalable inference backends with 99.9% uptime.

EDUCATION

Bachelor of Technology (B.Tech) in Computer Science & Engineering

State Technical University | CGPA: 8.8 / 10.0 | Graduated: 2024

TECHNICAL SKILLS

Languages & Frameworks:	Python, TypeScript, React, FastAPI, Node.js, PyTorch, SQL
Cloud & DevOps:	AWS (S3, Lambda, ECS), Docker, Kubernetes, CI/CD, PostgreSQL, Redis
AI & Specialized:	LLM Fine-tuning, RAG Architectures, Vector Embeddings, LangChain, vLLM

WORK EXPERIENCE & PROJECTS

AI Platform Engineer (Internship / Project) - Enterprise AI Lab

- Engineered an enterprise semantic search engine reducing document retrieval latency by 45%.
- Scaled async inference backend serving over 15000 daily API queries with 99.9% uptime.
- Built autonomous multi-agent copilot using LangChain and FastAPI boosting team velocity by 30%.

Autonomous Decision Intelligence Engine

- Designed and deployed an explainable decision intelligence pipeline handling 1000+ users.
- Integrated SHAP feature attribution to explain predictive factor weights in real time.
- Automated evaluation metrics and CI testing with 95% unit test coverage.

CERTIFICATIONS

- AWS Certified Solutions Architect - Associate
- DeepLearning.AI Generative AI & Large Language Models Specialization